

A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is a light green color. They are positioned diagonally, with the blue one in front of the green one.

Proof Positive

Can a machine find love on Twitter?



What did you want to know?

Public opinion about SXSW

What can we emphasize?

Do you want to know what went wrong?



Your data

~10,000 tweets were analyzed

The messages were graciously labeled and stored by an outside party:
<https://data.world/crowdflower/brands-and-product-emotions>



Methods

Natural Language Processing

“Count Vectorization”, aka the easy method, worked best





“Signs Point to Yes”


However, there were a couple key limitations

Your data = Proof of concept




“My reply is no”

The good news:
most people
liked SXSW



“Reply Hazy, Try Again”


Even identical messages were not always categorized together



Results

The model achieved 73% accuracy on unknown data.

“Positive” = only 30% of the data, but the model correctly labeled nearly 60% of those tweets.





How can we improve? For your team:

Size of the data pool

Tweets as a medium

More reliability in categorization



How can we improve? For you:

Not only will those raise our accuracy, but we may be able to hunt down what made people feel negatively about the experience.

Negative experiences hit us harder than positive ones do.



Potential Pitfalls

Resources spent labeling

Human error



Thank you

Feel free to reach me at:

Ferdinand.Beamer@gmail.com

or

www.linkedin.com/in/ferdinand-beaman

Credit to:

<https://www.cleanpng.com/png-magic-8-ball-8-ball-pool-eight-ball-crystal-ball-8-710716/>